

```
# To load from colab
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
import tensorflow as tf
print(tf.__version__)
```

2.4.0

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
from tensorflow.keras import layers
from tensorflow.keras import Model
from tensorflow.keras.applications import Xception
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import os
import matplotlib.pyplot as plt
import pandas as pd
from keras.callbacks import ReduceLROnPlateau
```

```
datagen = ImageDataGenerator( rescale = 1.0/255,
                              width_shift_range=0.1,
                              height_shift_range=0.1,
                              zoom_range=0.2,
                              vertical_flip=True,
                              horizontal_flip=True,
                              fill_mode='nearest')
```

#Load data

```
train_generator = datagen.flow_from_directory(
    '/content/drive/My Drive/Assgn-2/CV_Assignment3/intel-image-classification/seg_train/s',
    target_size=(150, 150),
    shuffle=True,
    batch_size=32,
    class_mode="categorical")
```

```
validation_generator = datagen.flow_from_directory(
    '/content/drive/My Drive/Assgn-2/CV_Assignment3/intel-image-classification/seg_train/s',
    target_size=(150, 150),
    batch_size=128,
    shuffle=True,
    class_mode="categorical")
```

```
test_generator = datagen.flow_from_directory(
    '/content/drive/My Drive/Assgn-2/CV_Assignment3/intel-image-classification/seg_test/se
```

```
target_size=(150, 150),
shuffle=True,
batch_size=128,
class_mode="categorical",
)
```

```
Found 14034 images belonging to 6 classes.
Found 14034 images belonging to 6 classes.
Found 3000 images belonging to 6 classes.
```

```
#load the pretrained model
```

```
from keras.applications.inception_v3 import InceptionV3
inceptionV3 = InceptionV3(include_top= False, input_shape=(150,150,3))
```

```
for layer in inceptionV3.layers:
    layer.trainable = False
```

```
last_layer = inceptionV3.get_layer('mixed9')
```

```
print('last layer output shape: ', last_layer.output_shape)
```

```
last_output = last_layer.output
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception\_v3/inception\_v3\_weights\_tf\_dim\_ordering\_tf\_data\_format.h5
87916544/87910968 [=====] - 1s 0us/step
last layer output shape: (None, 3, 3, 2048)
```



```
LearningRateScheduler = ReduceLROnPlateau(monitor='val_acc',
                                           patience=1,
                                           verbose=1,
                                           factor=0.20,
                                           min_lr=0.000001)
```

```
x = layers.Flatten()(last_output)
x = layers.Dense(1024, activation='relu')(x)
x = layers.Dropout(0.2)(x)
x = layers.Dense(6, activation='softmax')(x)
```

```
model = Model(inceptionV3.input, x)
```

```
model.compile(optimizer = tf.keras.optimizers.Adam(),
              loss = 'categorical_crossentropy',
              metrics = ['acc'])
```

```
model.summary()
```

conv2d_82 (Conv2D)	(None, 3, 3, 384)	442368	activation_81[0][0] ^
conv2d_83 (Conv2D)	(None, 3, 3, 384)	442368	activation_81[0][0]
average_pooling2d_7 (AveragePool)	(None, 3, 3, 1280)	0	mixed8[0][0]
conv2d_76 (Conv2D)	(None, 3, 3, 320)	409600	mixed8[0][0]

batch_normalization_78 (BatchNo	(None, 3, 3, 384)	1152	conv2d_78[0][0]
batch_normalization_79 (BatchNo	(None, 3, 3, 384)	1152	conv2d_79[0][0]
batch_normalization_82 (BatchNo	(None, 3, 3, 384)	1152	conv2d_82[0][0]
batch_normalization_83 (BatchNo	(None, 3, 3, 384)	1152	conv2d_83[0][0]
conv2d_84 (Conv2D)	(None, 3, 3, 192)	245760	average_pooling2d
batch_normalization_76 (BatchNo	(None, 3, 3, 320)	960	conv2d_76[0][0]
activation_78 (Activation)	(None, 3, 3, 384)	0	batch_normalizati
activation_79 (Activation)	(None, 3, 3, 384)	0	batch_normalizati
activation_82 (Activation)	(None, 3, 3, 384)	0	batch_normalizati
activation_83 (Activation)	(None, 3, 3, 384)	0	batch_normalizati
batch_normalization_84 (BatchNo	(None, 3, 3, 192)	576	conv2d_84[0][0]
activation_76 (Activation)	(None, 3, 3, 320)	0	batch_normalizati
mixed9_0 (Concatenate)	(None, 3, 3, 768)	0	activation_78[0][activation_79[0][
concatenate (Concatenate)	(None, 3, 3, 768)	0	activation_82[0][activation_83[0][
activation_84 (Activation)	(None, 3, 3, 192)	0	batch_normalizati
mixed9 (Concatenate)	(None, 3, 3, 2048)	0	activation_76[0][mixed9_0[0][0] concatenate[0][0] activation_84[0][
flatten (Flatten)	(None, 18432)	0	mixed9[0][0]
dense (Dense)	(None, 1024)	18875392	flatten[0][0]
dropout (Dropout)	(None, 1024)	0	dense[0][0]
dense_1 (Dense)	(None, 6)	6150	dropout[0][0]
=====			
Total params: 34,604,262			
Trainable params: 18,881,542			
Non-trainable params: 15,722,720			

```

history = model.fit(train_generator,
                    epochs = 10,
                    verbose = 1,
                    validation_data = validation_generator,
                    callbacks=[LearningRateScheduler])

```

```

Epoch 1/10
439/439 [=====] - 3436s 8s/step - loss: 0.6832 - acc: 0.7915
Epoch 2/10
439/439 [=====] - 187s 427ms/step - loss: 0.3216 - acc: 0.88
Epoch 3/10
439/439 [=====] - 187s 426ms/step - loss: 0.3278 - acc: 0.88

Epoch 00003: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
Epoch 4/10
439/439 [=====] - 187s 427ms/step - loss: 0.2561 - acc: 0.90
Epoch 5/10
439/439 [=====] - 186s 425ms/step - loss: 0.2289 - acc: 0.91
Epoch 6/10
439/439 [=====] - 188s 428ms/step - loss: 0.2176 - acc: 0.92
Epoch 7/10
439/439 [=====] - 187s 427ms/step - loss: 0.2150 - acc: 0.91
Epoch 8/10
439/439 [=====] - 184s 421ms/step - loss: 0.2033 - acc: 0.92

Epoch 00008: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.
Epoch 9/10
439/439 [=====] - 185s 421ms/step - loss: 0.1944 - acc: 0.92
Epoch 10/10
439/439 [=====] - 184s 420ms/step - loss: 0.1886 - acc: 0.93

```

Predictions on test dataset

```

model.save_weights('/content/drive/My Drive/Assgn-2/CV_Assignment3')
test_generator.reset()
pred=model.predict_generator(test_generator,
verbose=1)

/usr/local/lib/python3.6/dist-packages/tensorflow/python/keras/engine/training.py:196
warnings.warn("`Model.predict_generator` is deprecated and '
24/24 [=====] - 951s 41s/step

```

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score, accuracy_score
ypred=predicted_class_indices
ytest=test_generator.labels
confusion_mat=tf.math.confusion_matrix(
    ytest,ypred,num_classes=None, weights=None, dtype=tf.dtypes.int32,
    name=None
)

acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))

```

```

plt.plot(epochs, acc, 'r', label='Training accuracy')

```


```

plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()

plt.plot(epochs, loss, 'r', label='Training Loss')
plt.plot(epochs, val_loss, 'b', label='Validation Loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

```



```

import numpy as np
from keras.preprocessing import image

file_path = '/content/drive/My Drive/Assgn-2/CV_Assignment3/intel-image-classification/se

test_image = image.load_img(file_path, target_size = (150, 150))
plt.imshow(test_image)
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = model.predict(test_image)

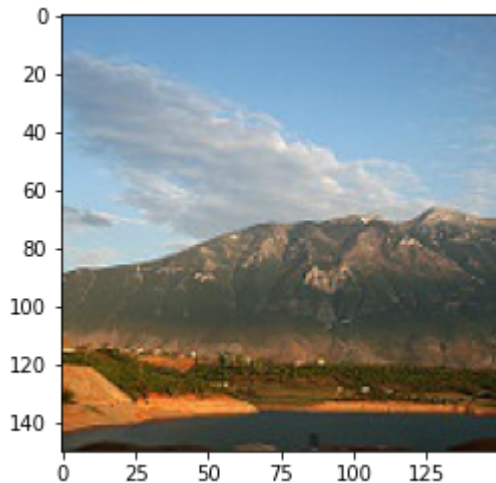
classes=train_generator.class_indices

```

```
classes=['buildings','forest','glacier','mountain','sea','street']
```

```
print("Prediction for this image is :",classes[np.argmax(result)])
```

Prediction for this image is : sea



```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score, accuracy_score
ypred=predicted_class_indices
ytest=test_generator.labels
confusion_mat=tf.math.confusion_matrix(
    ytest,ypred,num_classes=None, weights=None, dtype=tf.dtypes.int32,
    name=None
)
print('F1 Score =', f1_score(ytest, ypred,average='micro'))
print('Accuracy =', accuracy_score(ytest, ypred))
FP = confusion_mat[0][1]
TP = confusion_mat[1][1]
print('True Positive =', TP)
print('False Positive =', FP)

F1 Score = 0.17333333333333334
Accuracy = 0.17333333333333334
True Positive = tf.Tensor(71, shape=(), dtype=int32)
False Positive = tf.Tensor(75, shape=(), dtype=int32)
```

